

How to Read a Research Paper

Nick Feamster and Alex Gray

CS 7001

August 30, 2006

Administrivia

- Check web page for lecture schedule/assignment writeup
- Use class time to meet as a group
 - Simplifies scheduling: everyone should be free
 - Meet wherever you like
 - We will be at the classroom on Friday
- Subscribe to mailing lists if you haven't done so
 - cs7001

Agenda

- Finding and crafting good problems (conclude)
- Reading research papers
 - Should help with PS1

How do you find ideas/problems?

- Frustrations
 - Your own
 - Others'... how to find out about these?
- Read, read, read to stay relevant
 - Mailing lists
 - Conferences (“real world” ones, in particular)
 - Certain news rags (*Economist*, *Tech Review*, etc.)
 - Can often be a good source of ideas for applying one problem domain to another, etc.

Where *not* to find ideas/problems

- Your imagination
 - Don't make problems up. Plenty of real ones out there...
 - The new approach, question, etc., should come from your thoughts. The problem itself should not.
 - Asking for feedback is key (e.g., asking people to read your papers...today's topic)
- Conference proceedings (sometimes)
 - The last 80% of papers on a topic solve the last 5% of the problem
 - Often, that 5% “doesn't matter”
 - Engineering or time travel will solve it...

Making a good problem from a big one

- Baby steps: simplifications of the problem
 - Taking care to keep assumptions valid
- Special cases of the problem
 - When you can't find the answer, sometimes changing the question works.
- Multiple perspectives on the same problem
- Surrounding the problem by knowing all aspects; massive knowledge
- Big analogical leaps

Questions

- Why read research papers?
- What to read?
- How to read?
 - Quickly
 - Selectively
 - Critically

Why Read Research Papers?

- Keep up with the field
- Understand different perspectives on your problem
- Learn the lingo
- Avoid trying to reinvent the wheel
- Have you picked a problem of the right size
- Learn writing style and culture
- Pick problems
- Combine problems

Why Read Research Papers?

- Look for new ideas
- Understand the context of your own work
 - General problem area
 - Related results
- Build your body of knowledge
- Keep up with a research area
- Someone asks you for feedback

Where to start? So much to read!

- Start with the classics
 - These themes will recur
 - Can provide context for the current readings
 - These folks often tend to have written books (often have a good writing style, etc.)
 - Where to find them? Course notes, Springer notes (?), area exam lists, your own hunting, test of time awards, etc.
- Check the references
- Conference proceedings!

What to Read

- Proceedings of top conferences
 - Skim everything
 - Read the subset that directly applies
 - Each proceedings will probably have < 3 papers that you will read-end to end
- The lost gems
 - In your area: Follow the people you respect
 - Develop rapport with more senior mavens
 - Out of your area: You'll hear about them
 - No need to be the maven out of your area
 - (What do to when your paper becomes a lost gem?)
- Journals: Importance is variable by field. Find out!

How to Quickly Understand a Paper?

- Abstract, Introduction and conclusions first!
- Section headings
- Pictures, graphs, etc.
- Explanation of previous work
- See if you can summarize the paper in a sentence or two
- Check out the authors to get context

Takeaway: Read Selectively

- Don't read every paper the same way
 - Not necessary (or advisable) to read every paper end-to-end
 - Different from *reviewing* papers
- Use your time wisely
 - Time management task!
 - Too much to read: Can't even read all of the work related to yours
 - How to decide which papers to read?

Reading Selectively

- Like peeling an onion: many layers of detail
 - Go breadth first
 - How much detail you need depends on what you're going to do with the work
- One approach to reading a paper
 - Read abstract, intro, and conclusions first
 - At this point, decide whether you need to read on
 - Ask the questions about “Is this a good problem/idea?”
 - Need to understand results better
 - Read experiments/results last
 - Will take 80% of the time
 - Often, summary of results in intro is enough

Reading Selectively (cont.)

- Speed Reading
 - Only read first and last sentence of each paragraph
- Can quickly get a sense for things like formality level
- Evaluation sections
 - Don't read unless:
 - You need the detail
 - The technique itself is interesting
 - Look at graphs and captions first
- Reviewing: focus on flaws

Answer key questions first

- Is the main point obvious?
 - What is the key “intellectual nugget”?
- Are there any experimental results?
 - Does the rest of the paper have anything to do with abstract and intro? (i.e., are the claims backed up?)
- Are there any theorems?
- Did they compare to any competition?

Reading Critically

- Multi-pass approach
 - Read with a pen
 - Return later to parts you don't understand
 - Write down main points at the top of 1st page (quick reference)
- Summarize the paper in 1-2 sentences
 - If you can't do this after reading the abstract/intro/conclusion, think hard about whether you want to read on
 - Useful when you review a paper, too (signal)
- Ask these questions
 - Does the paper address a “good” problem?
 - Does the paper follow through on the claims?
 - What assumptions were made? (valid?)
 - Soundness of methods/experimentation

Reading a Paper End-to-End

- Classic papers
- Papers that are directly related to your research

When Someone Asks You for Comments...

- Personal favors are important
 - You will need readers for your own drafts
 - Feedback is critical to research
 - Still, you've got stuff to do. Can't read every word!
- It is important to prioritize
 - Your advisor, friends, and immediate colleagues
 - People doing work directly related to yours ("contemporaries")
 - Sometimes can give you a sneak peak into what others are up to
 - Work by colleagues unrelated to your research
- High-level feedback can go a long way