

# **PS 1 Lessons**

## **(Plus: Reading vs. Reviewing)**

Nick Feamster and Alex Gray

CS 7001

September 25, 2006

# Networking Group Open House

- This Wednesday, 4-7 p.m.
- MiRC 102 (across from the entrance of CoC)
- Content
  - Summary of the networking research at GaTech
    - Meet the faculty
  - Presentation from Abhishek Kumar
  - Student posters
- Plenty of food and drink

# Agenda

- Administrivia
  - PS 2 deadline extended two weeks
  - More time to come up with great ideas
- Today: Lessons and feedback from PS 1
- Reading vs. Reviewing
  - What goes in a review?
  - How to write for a reviewer?

# Lessons on Reading

- What did you learn?
- Did you read/understand every paper?
  - What did you do about the ones you didn't understand?
- Did you re-read any papers? Why or why not?
- How did you divide the work?
- Suggestions for the assignment?

# Lessons: Reading

# Some Observations: Reading

- Impossible to read the entire proceedings
  - Important to identify good ideas quickly
  - Key: Division of work and help from colleagues
- Questions the writeups didn't always address
  - Is the problem being demonstrated real?
    - Where can you find out? Evidence?
  - Are the claims of the introduction backed up?
    - Is there an evaluation? How credible is it?
  - What is the context of the research?
    - Open new area/perspective
    - “Close the book” on a longstanding problem

# Selection: Similar to Review Process

- If you divided the work of reading, shared notes on papers, and argued to consensus...
  - Program committees take a similar approach
- Consider the lessons when thinking about how to present your own research
  - Need to make the reader's job easier
    - People have limited time and patience
    - Your paper is being read late at night by someone who is interested in finishing the task quickly

# Reading vs. Reviewing



# Reading

- Primary goal: Knowledge transfer
  - Details needed for your own research
    - Front-to-back read is often required
  - Building the arsenal
    - Knowledge of related tools/techniques
    - Keeping up with what's “hot”
    - Knowing what problems are “solved”
- 80% - 90% of papers fall into the second category
  - $40 \text{ papers/conf/year} * 5 \text{ years} * 3 \text{ conf} = 600 \text{ papers}$
  - Your thesis: 100 citations

# Reviewing

- Goal: Determine whether the paper merits publication in the venue where it was submitted
  - Must consider appropriateness for venue/audience
  - Different venues have different standards, expectations
  - Often, a front-to-back read is required
- Secondary: Help the authors improve their paper
  - The end goal is not a binary decision
    - If you saw the paper a second time (in submission, or the final version), what needs to be fixed
  - Don't spend longer reviewing the paper than the authors did writing it

# What a review should include

# What a review should include

- A clear, concise summary of the paper
  - Signals to the authors that you read and understood the main point of the paper
  - If you didn't “get it”, helps them understand why
- A clear recommendation
  - Accept, reject, borderline
- High-level, technical comments
- Detailed comments
  - Punctuation, spelling, unclear sentences, etc.
  - Don't go overboard if the paper is very bad

# Lessons: Writing Your Own Papers

- Properties of the papers you selected?
  - Contributions clear?
  - Shortcomings acknowledged or addressed?
  - Thoughts on writing style?
    - Declarative or passive tone?
- Did you see approaches you wouldn't take?

# **Lessons for Writing**

## **(More on this in Later Lecture)**

# From the Reviewer's Point of View

## *A Reviewer's Plea* by Mark Allman



I have just completed my first (maybe only!) stint on the program committee for ACM SIGCOMM's annual symposium. ... This experience was eye-opening in that I now realize that the community generates a large amount of **\*\*junk\*\***.

And, I don't mean "junk" in the sense of bad ideas. In general, each paper I reviewed for SIGCOMM this year had at least a nugget of an interesting idea. I mean "junk" in the "it took me twice as long to read this as it should have" sense. In other words, the paper was sloppy.

I had a very difficult time trying to figure out what the paper was proposing or what the experiments were really showing. The following is a short list of **\*easy\*** ways to improve your research papers.

# Point #1: Spelling and Proofreading

- If you don't care about your paper, why should the reviewer?
- If you are sloppy with your text, what should the reviewer think about your experiments?
- “Sloppy papers indicate sloppy research”
- Your paper must be understood to be accepted
  - Note the difference from reading, here, where often you must understand a paper that's been accepted



# Point #2: Background Material

- Consider the audience when presenting background material
- Don't waste time presenting concepts that 90% of readers are likely to know
- At the same time, your coverage of related work must be complete
  - Missing a paper in a major conference or a highly cited paper is not excusable
  - Make certain to cite related work by the PC :)

# Point #3: Make It Easy on the Reviewer

- Design plots and diagrams for maximum information transfer
  - Future lecture on presenting information visually
  - Don't forget to label axes clearly, use units, etc., to make it clear what is being evaluated
- Clear, readable fonts and typography help

# Suggestions for PS 2

- Perform a literature review to determine whether your problem/approach is important, new, etc.
  - We expect to see references
- Consider how you would *evaluate* your idea
  - Proofs/analysis
  - Simulation
  - Implementation/Real-world deployment